

# SE6103 Parallel and Distributed Systems

## Assignment 3

19APSE4325 – Tharindu Lakmal

### Question 2:

```
1 version: '3'
2 services:
3   namenode:
4     image: bde2020/hadoop-namenode:latest
5     container_name: namenode
6     environment:
7       - CLUSTER_NAME=testhadoop
8       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
9     ports:
10      - "9870:9870" # Web UI
```

```
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker-compose up -d
time="2025-01-06T14:41:44+05:30" level=warning msg="D:\xml\ParallelDistributed\Hadoop_11_25\docker-compose.yml: 'version' is obsolete"
[+] Running 1/2
 ✓ Network hadoop_11_25_default Created
   Container namenode
Error response from daemon: Conflict. The container name "/namenode" is already in use by container "8b35bd76bd648953015cbb9ef7f54665810118a2248bcb6d73e7c8fdd62ab7". You
have to remove (or rename) that container to be able to reuse that name.
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -mkdir -p /input

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put sample.txt /input
put: 'sample.txt': No such file or directory

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put ./sample.txt /input
put: './sample.txt': No such file or directory

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put ./sample.txt /input
put: './sample.txt': No such file or directory

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /input
```

```
1 version: '3'
2 services:
3   namenode:
4     image: bde2020/hadoop-namenode:latest
5     container_name: namenode
6     environment:
7       - CLUSTER_NAME=testhadoop
8       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
9     ports:
10      - "9870:9870" # Web UI
```

```
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /input

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker cp ./sample.txt namenode:/sample.txt
Successfully copied 2.05kB to namenode:/sample.txt
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -put ./sample.txt /input
2025-01-06 09:20:04,764 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false

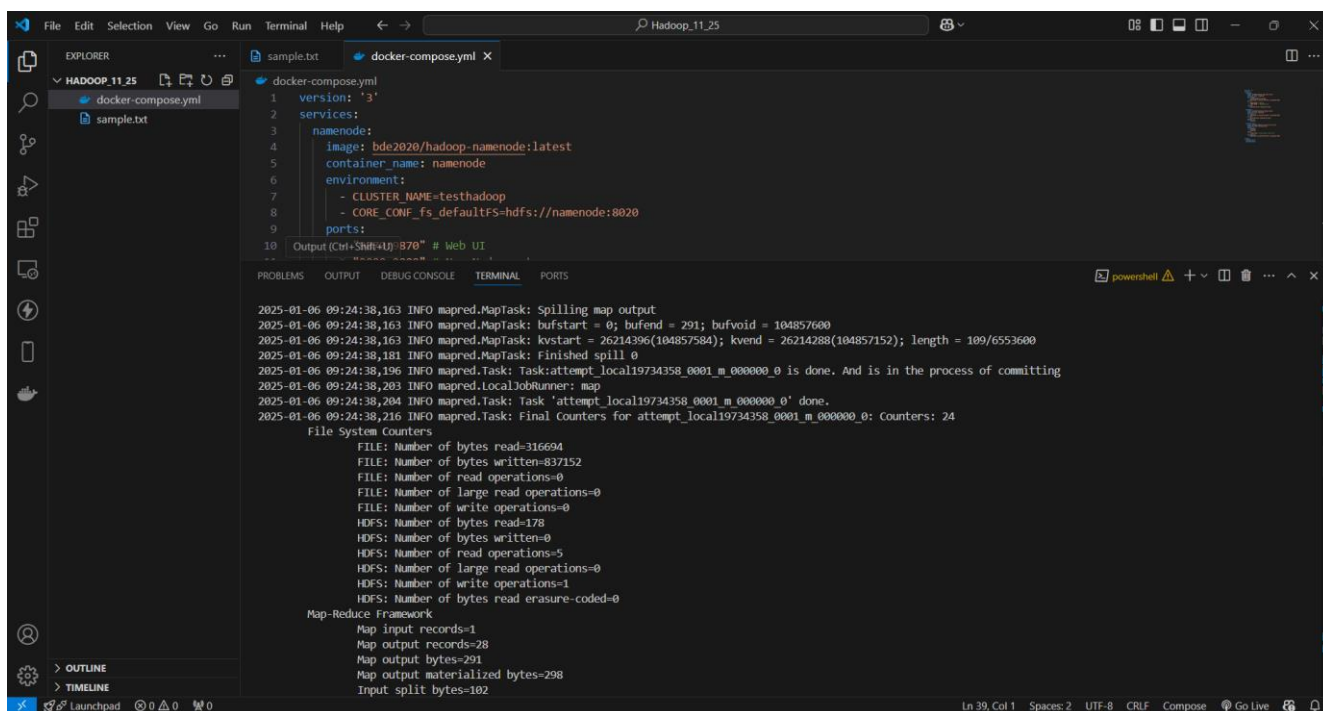
What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /input
Found 1 items
-rw-r--r-- 3 root supergroup 178 2025-01-06 09:20 /input/sample.txt

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
```

```
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hadoop jar /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar wordcount /input /output
JAR does not exist or is not a normal file: /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-*.jar

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hadoop jar /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.2.jar wordcount /input /ou
tput
JAR does not exist or is not a normal file: /opt/hadoop/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.2.jar

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image + docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar wordcount /inp
ut /output
2025-01-06 09:24:36,253 INFO impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2025-01-06 09:24:36,437 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2025-01-06 09:24:36,437 INFO impl.MetricsSystemImpl: JobTracker metrics system started
2025-01-06 09:24:36,799 INFO input.FileInputFormat: Total input files to process : 1
2025-01-06 09:24:36,881 INFO mapreduce.JobSubmitter: number of splits:1
2025-01-06 09:24:37,460 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local19734358_0001
2025-01-06 09:24:37,460 INFO mapreduce.JobSubmitter: Executing with tokens: []
2025-01-06 09:24:37,601 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2025-01-06 09:24:37,602 INFO mapreduce.Job: Running job: job_local19734358_0001
2025-01-06 09:24:37,604 INFO mapred.LocalJobRunner: OutputCommitter set in config null
```

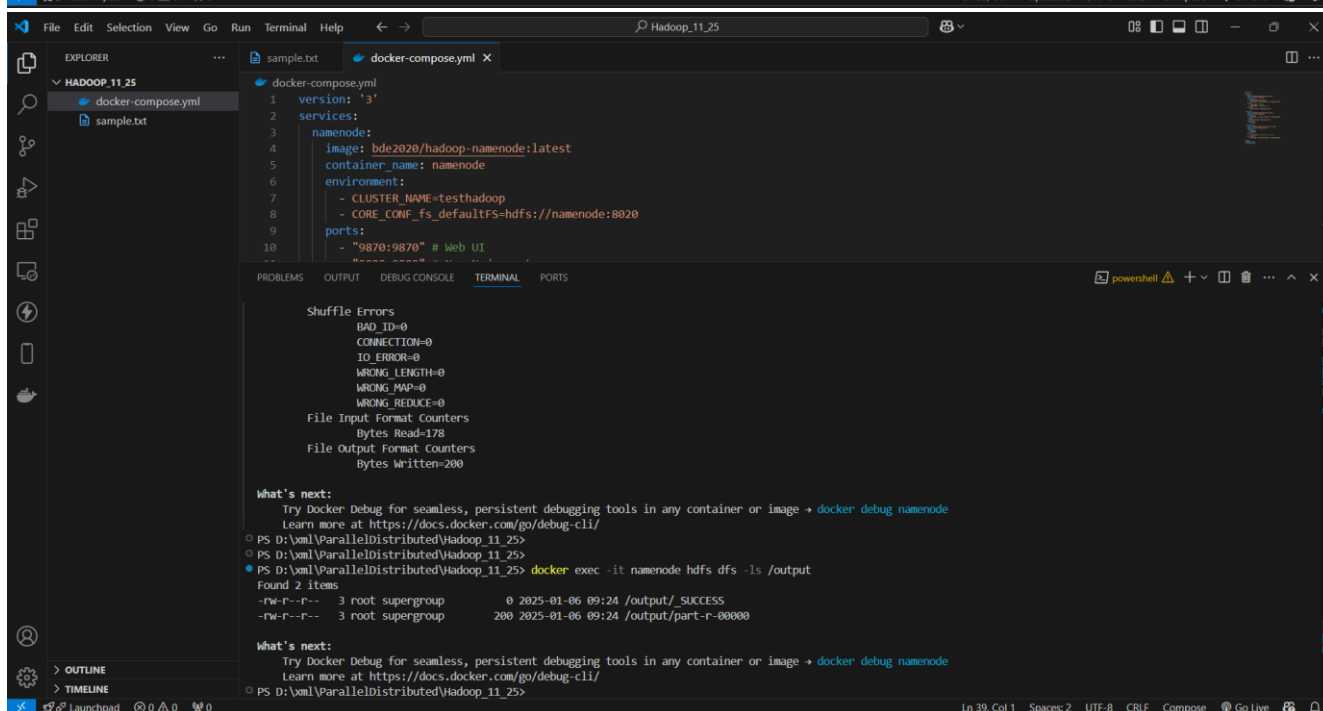


The screenshot shows the Visual Studio Code interface with a file explorer on the left containing 'sample.txt' and 'docker-compose.yml'. The main editor displays the 'docker-compose.yml' file with the following content:

```
1 version: '3'
2 services:
3   namenode:
4     image: bde2020/hadoop-namenode:latest
5     container_name: namenode
6     environment:
7       - CLUSTER_NAME=testhadoop
8       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
9     ports:
10      - "9870:9870" # Web UI
```

The terminal window at the bottom shows Hadoop logs for a namenode task. The logs include information about map tasks, file system counters, and map-reduce framework statistics.

```
2025-01-06 09:24:38,163 INFO mapred.MapTask: Spilling map output
2025-01-06 09:24:38,163 INFO mapred.MapTask: bufstart = 0; bufend = 291; bufvoid = 104857600
2025-01-06 09:24:38,163 INFO mapred.MapTask: kvstart = 26214396(104857584); kvend = 26214288(104857152); length = 109/6553600
2025-01-06 09:24:38,196 INFO mapred.Task: Task:attempt_local19734358_0001_m_000000_0 is done. And is in the process of committing
2025-01-06 09:24:38,203 INFO mapred.LocalJobRunner: map
2025-01-06 09:24:38,204 INFO mapred.Task: Task 'attempt_local19734358_0001_m_000000_0' done.
2025-01-06 09:24:38,216 INFO mapred.Task: Final Counters for attempt_local19734358_0001_m_000000_0: Counters: 24
  File System Counters
    FILE: Number of bytes read=316694
    FILE: Number of bytes written=837152
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=178
    HDFS: Number of bytes written=0
    HDFS: Number of read operations=5
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=1
    HDFS: Number of bytes read erasure-coded=0
  Map-Reduce Framework
    Map input records=1
    Map output records=28
    Map output bytes=201
    Map output materialized bytes=298
    Input split bytes=102
```



The screenshot shows the Visual Studio Code interface with the same file explorer. The main editor displays the 'docker-compose.yml' file with the following content:

```
1 version: '3'
2 services:
3   namenode:
4     image: bde2020/hadoop-namenode:latest
5     container_name: namenode
6     environment:
7       - CLUSTER_NAME=testhadoop
8       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
9     ports:
10      - "9870:9870" # Web UI
```

The terminal window at the bottom shows the output of a Docker command. The output includes a list of shuffle errors, file input/output format counters, and a list of what's next steps.

```
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=178
File Output Format Counters
  Bytes Written=200

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25>
PS D:\xml\ParallelDistributed\Hadoop_11_25> docker exec -it namenode hdfs dfs -ls /output
Found 2 items
-rw-r--r--   3 root supergroup          0 2025-01-06 09:24 /output/_SUCCESS
-rw-r--r--   3 root supergroup    200 2025-01-06 09:24 /output/part-r-000000

What's next:
Try Docker Debug for seamless, persistent debugging tools in any container or image → docker debug namenode
Learn more at https://docs.docker.com/go/debug-cli/
PS D:\xml\ParallelDistributed\Hadoop_11_25>
```

File Edit Selection View Go Run Terminal Help Hadoop\_11\_25

EXPLORER

- HADOOP\_11\_25
  - docker-compose.yml
  - sample.txt

sample.txt

```
1 version: '3'
2 services:
3   namenode:
4     image: bde2020/hadoop-namenode:latest
5     container_name: namenode
6     environment:
7       - CLUSTER_NAME=testhadoop
8       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
9     ports:
10      - "9870:9870" # Web UI
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

PS D:\xaml\ParallelDistributed\Hadoop\_11\_25>
PS D:\xaml\ParallelDistributed\Hadoop\_11\_25> docker exec -it namenode hdfs dfs -cat /output/part-r-00000
2025-01-06 09:26:15,885 INFO sasl.SaslDataTransferClient: SASL encryption trust check: localhostTrusted = false, remoteHostTrusted = false
Hadoop 2
about 1
across 1
amounts 1
an 1
and 1
computers. 1
data 1
framework 2
is 2
large 1
learn 1
network 1
networking 1
of 2
open-source 1
processes 1
stores 1
that 1
to 1
very 1

Launchpad 0 0 0 0

Ln 39, Col 1 Spaces: 2 UTF-8 CRLF Compose Go Live

File Edit Selection View Go Run Terminal Help Hadoop\_11\_25

EXPLORER

- HADOOP\_11\_25
  - docker-compose.yml
  - sample.txt

sample.txt

```
1 version: '3'
2 services:
3   namenode:
4     image: bde2020/hadoop-namenode:latest
5     container_name: namenode
6     environment:
7       - CLUSTER_NAME=testhadoop
8       - CORE_CONF_fs_defaultFS=hdfs://namenode:8020
9     ports:
10      - "9870:9870" # Web UI
```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

is 2
large 1
learn 1
network 1
networking 1
of 2
open-source 1
processes 1
stores 1
that 1
to 1
very 1

What's next:  
Try Docker Debug for seamless, persistent debugging tools in any container or image + `docker debug namenode`  
Learn more at <https://docs.docker.com/go/debug-cli/>

PS D:\xaml\ParallelDistributed\Hadoop\_11\_25>  
PS D:\xaml\ParallelDistributed\Hadoop\_11\_25>  
PS D:\xaml\ParallelDistributed\Hadoop\_11\_25> docker-compose down  
time="2025-01-06T14:56:35+05:30" level=warning msg="D:\xaml\ParallelDistributed\Hadoop\_11\_25\docker-compose.yml: 'version' is obsolete"  
[+] Running 1/1  
✓ Network hadoop\_11\_25\_default Removed  
PS D:\xaml\ParallelDistributed\Hadoop\_11\_25>

Launchpad 0 0 0 0

Ln 39, Col 1 Spaces: 2 UTF-8 CRLF Compose Go Live